

Mahendra Kamireddy

Data Engineer

Meetmahendra204@gmail.com | +91-6302840590 | [Github](#) | [Medium](#)

Professional Summary

Data Engineer with 3+ years of experience in building, optimizing data pipelines and enhancing reporting workflows using AWS Services (S3, Athena, Glue, Redshift, RDS, Lambda, EMR), Microsoft Azure Services (ADLS, ADF, Azure Databricks, Azure SQL DB), GCP Services (GCS, Dataflow, Big Query, Dataproc), Snowflake, and Data Manipulations, analysis, and Validations using SQL, Python, PySpark, along with building dashboards using Power BI, in delivering strategic insights and operational efficiency improvements to stakeholders.

Technical Skills

Programming Languages: SQL, Python

Data Bases: MS SQL Server, MySQL, PostgreSQL

Cloud Technologies: AWS (S3, Athena, Glue, Redshift, EMR, Lambda, IAM, Databricks),
Microsoft Azure (ADF, ADLS, Azure SQLDB, Synapse Analytics), Microsoft Fabric,
GCP (GCS, Dataflow, Big Query, Dataproc), Snowflake.

Miscellaneous: Power BI, Data Modelling, DAX, Pandas, Numpy, Spark, PySpark, Airflow, Machine Learning, Statistics, Data Analysis, MS Excel, Git, JIRA.

Work Experience

Persistent Systems Limited

Jan 2026 - Present

Client: Wells Fargo Data Engineer / Lead Software Engineer

- Executed secure migration of 10+ TB of sensitive data from on-premise SQL Server, Oracle, and MongoDB systems to Google Cloud Platform (GCP) via OpenShift Container Platform (OCP), ensuring encryption in transit and at rest with zero data loss.
- Designed and implemented full + incremental load strategies, reducing refresh cycle times and improving pipeline efficiency for daily and weekly security data ingestion.
- Automated ETL workflows with schema validation and audit logging, cutting manual intervention by 60% and strengthening compliance with Wells Fargo's enterprise security standards.
- Optimized GCP storage and compute costs through partitioning and incremental refresh logic, achieving cost savings compared to baseline migration estimates.
- Built monitoring dashboards and SLA tracking mechanisms, improving incident response time by 35% and ensuring proactive detection of pipeline failures.

Upwork (Freelance)

May 2024 – Jan 2026

Data Engineer / Data Analyst

- Building Scalable Data pipelines using Azure Services (ADF, ADLS Gen2, Azure Databricks, Azure DB, IAM, etc.), and AWS Services (Glue, S3, Redshift, Athena, Lambda, AWS EMR, AWS Databricks).
- Working on data analysis and reporting the results to the clients to take actions based on data and data analysis. Collaborated with clients to understand the requirements and design the most robust solutions to solve business problems.

e-Global Infotech

March 2024 – May 2024

(Client: Imperial Brands PLC)

Azure Data Engineer / Data Analyst

- Integrated data from various sources using Azure Data Factory and stored it in Azure Storage services like ADLS Gen2 and Azure SQL DB, ensuring that by building optimized data pipelines.
- Optimized an Azure Databricks job that was taking 1 hour to run, reducing the runtime to just 40 minutes by repartitioning the data and optimizing the spark configuration.
- Increased data storage efficiency by 20% through effective Azure Storage Services ADLS Gen2 by partitioning and compression.

Infosys Limited, Bangalore

June 2022 – Feb 2023

Big Data Engineer /System Engineer

- Developed scalable data pipelines for processing large datasets across multiple domains. Conducted exploratory data analysis (EDA) using Python, SQL, and Power BI, enhancing business insights.
- Designed interactive dashboards in Power BI to visualize complex financial and operational trends.
- Collaboration with researchers and analysts to structure datasets for AI/ML experimentation.

Projects

1. Tobacco Supply Chain Analytics Technologies: Azure Services – ADF, ADLS Gen2, Azure Databricks

- Designed an end-to-end supply chain analytics solution to monitor inventory turnover, supplier lead times, logistics costs, and stock availability for a manufacturing client.
- Integrated data from ERP systems, warehouse management systems, and logistics sources into a centralized Azure based data platform.
- Built robust ETL pipelines using Azure Data Factory and Microsoft Fabric to ingest, clean, and transform raw data.
- Modelled data using star schema (fact and dimension tables) to optimize query performance and reporting efficiency. Developed interactive Power BI dashboards with drill-downs, slicers, and KPI indicators for executive and operations teams.
- Implemented DAX measures for inventory aging, fill rate, supplier performance scoring, and cost variance analysis. Enabled automated data refresh and monitoring, ensuring near real-time visibility into supply chain operations.

2. Patient Analytics Infrastructure

Technologies: Azure, Power BI, SQL, Excel

- Designed and implemented a healthcare analytics platform to analyze patient outcomes, readmission rates, and treatment effectiveness.
- Integrated data from multiple healthcare sources including EMR systems, laboratory data, and insurance claims.
- Built scalable ETL workflows using Azure Data Factory to process and standardize healthcare datasets.
- Developed optimized SQL queries and data models to support clinical and operational reporting.
- Created interactive Power BI dashboards for hospital administrators and clinical leadership.

Education

PBR Visvodaya Institute of Technology & Science, Kavali, A.P, JNTUA
M.TECH. in Power Electronics, Electrical and Electronics Engineering.

Jan 2021

Course Works: Power Electronics, Machine Learning, Data Science and Analytics.